

Problem_3

January 20, 2023

1 Imports

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import os
import pandas as pd
from scipy.optimize import minimize

plt.rcParams['figure.figsize'] = (10.0, 7.0)
plt.rcParams["font.size"] = 16
plt.rcParams["font.family"] = "Serif"
```

```
[ ]: DATA_DIR = 'data'
SAVE_DIR = "plots"
```

2 The Problems

Here are the problems:

Apply the logistic regression (linear classifier) algorithm discussed in the lab session to predict next-day rain based on the 10 years of daily weather observations from many locations within a country. The dataset contains many factors taken into consideration to specify whether it rained or not on that particular day. The training and testing dataset is provided in the files titled 'weather_train.csv' and 'weather_test.csv', respectively. Carry out the following tasks as assignment problems:

1. Inspect and plot some portion of the training data using pandas. Segregate the training and testing data into two separate variables consisting of 'feature values' and corresponding 'predictions' (the prediction column is titled 'RainTomorrow' in the dataset). To simplify the problem a bit, clean the whole data by carrying out the following sub-tasks:
 - Convert the predictions in the binary format by using '1' for 'YES' and '0' for 'NO'.
 - Identify and drop the feature columns having datatype 'object'.
 - Identify cells having 'NaN' or 'NA' values and replace them with mean values of their respective columns.
 - Normalize all the feature values by scaling them between 0 and 1. The values in feature

matrix X can be normalized as:

$$X_{norm} = \frac{X - \min(X)}{\max(X) - \min(X)}$$

Execute the above sub-tasks and display some portion of the data and its head after each data cleaning step.

2. Classify the cleaned dataset using binary classification algorithm discussed in the class and calculate the optimized weights and training set accuracy for the model (use Truncated Newton's Method in SciPy for optimization).
3. Plot the cost history (J) vs. the number of iterations
4. Apply the trained model on the cleaned test dataset to predict the testing accuracy of the model.

2.1 The Approach

We can formulate the problem of logistic regression as an optimization problem. For this, let's define some quantities.

Suppose the input data is $X \in \mathbb{R}^{m \times n}$ where m is number of samples and n number of features. The target variable $\mathbf{y} \in \mathbb{R}^m$ has just two possible values, say 0 and 1.

2.1.1 Hypothesis Function

The hypothesis function for logistic regression is:

$$h(\theta) = z(\mathbf{w}^T X)$$

Here $\mathbf{w} \in \mathbb{R}^{n+1}$ is the weight including the bias term, that is $\mathbf{w} = [w_0, w_1, \dots, w_n]$ and z is a function, called the sigmoid function defined as:

$$z = \frac{1}{1 + e^{-x}}$$

2.1.2 Cost Function

The cost function which we will use for the logistic regression is

$$J(\hat{y}, \mathbf{y}) = \begin{cases} -\log(\hat{y}) & \text{if } \mathbf{y} = 1 \\ -\log(1 - \hat{y}) & \text{if } \mathbf{y} = 0 \end{cases}$$

This can be rewritten as:

$$J(\hat{y}, \mathbf{y}) = -\mathbf{y} \log(\hat{y}) - (1 - \mathbf{y}) \log(1 - \hat{y})$$

Note that J is function of \hat{y} which in turn is function of \mathbf{w} and X , so, J is function of $J(\mathbf{w}, X, \mathbf{y})$.

2.1.3 Optimization Problem

Using these definitions, we can formulate the optimization problem as:

$$\text{minimize}_{\mathbf{w}} J(\mathbf{w}, X, \mathbf{y})$$

3 Solving The Third Problem

3.1 Problem 3.1

```
[ ]: train = pd.read_csv(os.path.join(DATA_DIR, "weather_train.csv"))
test = pd.read_csv(os.path.join(DATA_DIR, "weather_test.csv"))
train.head()
```

```
[ ]:
      Date Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  \
0  01-12-2008  Albury    13.4    22.9      0.6          NaN        NaN
1  02-12-2008  Albury     7.4    25.1      0.0          NaN        NaN
2  03-12-2008  Albury    12.9    25.7      0.0          NaN        NaN
3  04-12-2008  Albury     9.2    28.0      0.0          NaN        NaN
4  05-12-2008  Albury    17.5    32.3      1.0          NaN        NaN

      WindGustDir  WindGustSpeed  WindDir9am  ...  Humidity9am  Humidity3pm  \
0              W             44.0          W  ...        71.0         22.0
1            WNW             44.0        NNW  ...        44.0         25.0
2            WSW             46.0          W  ...        38.0         30.0
3             NE             24.0          SE  ...        45.0         16.0
4              W             41.0        ENE  ...        82.0         33.0

      Pressure9am  Pressure3pm  Cloud9am  Cloud3pm  Temp9am  Temp3pm  RainToday  \
0         1007.7         1007.1        8.0      NaN     16.9     21.8         No
1         1010.6         1007.8        NaN      NaN     17.2     24.3         No
2         1007.6         1008.7        NaN      2.0     21.0     23.2         No
3         1017.6         1012.8        NaN      NaN     18.1     26.5         No
4         1010.8         1006.0        7.0      8.0     17.8     29.7         No

      RainTomorrow
0              No
1              No
2              No
3              No
4              No

[5 rows x 23 columns]
```

```
[ ]: test.head()
```

```
[ ]:
      Date      Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  \
0  11-01-2014  PearceRAAF    19.0    44.5      0.0          NaN        13.1
1  12-01-2014  PearceRAAF    31.2    44.3      0.0          NaN        12.2
2  13-01-2014  PearceRAAF    19.3    30.3      0.0          NaN         9.4
3  14-01-2014  PearceRAAF    14.0    29.7      0.0          NaN        13.1
4  15-01-2014  PearceRAAF    12.7    27.9      0.0          NaN        12.4

      WindGustDir  WindGustSpeed  WindDir9am  ...  Humidity9am  Humidity3pm  \
```

0	E	54.0	E ...	27.0	10.0
1	E	54.0	N ...	10.0	22.0
2	WSW	46.0	SSW ...	63.0	43.0
3	WSW	44.0	SSE ...	43.0	29.0
4	W	50.0	NW ...	48.0	34.0

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	\
0	1015.3	1009.7	NaN	NaN	31.2	42.9	No	
1	1007.7	1007.8	NaN	NaN	40.2	35.8	No	
2	1010.9	1009.5	7.0	2.0	23.2	27.6	No	
3	1012.2	1009.5	0.0	0.0	23.0	29.0	No	
4	1008.4	1008.4	1.0	3.0	23.4	25.7	No	

	RainTomorrow
0	No
1	No
2	No
3	No
4	No

[5 rows x 23 columns]

```
[ ]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 52062 entries, 0 to 52061
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  52062 non-null  object
1   Location              52062 non-null  object
2   MinTemp              51538 non-null  float64
3   MaxTemp              51672 non-null  float64
4   Rainfall             50766 non-null  float64
5   Evaporation          24047 non-null  float64
6   Sunshine             18441 non-null  float64
7   WindGustDir          46533 non-null  object
8   WindGustSpeed        46540 non-null  float64
9   WindDir9am           46068 non-null  object
10  WindDir3pm           49611 non-null  object
11  WindSpeed9am         50930 non-null  float64
12  WindSpeed3pm         50306 non-null  float64
13  Humidity9am          51272 non-null  float64
14  Humidity3pm          50667 non-null  float64
15  Pressure9am          45067 non-null  float64
16  Pressure3pm          45117 non-null  float64
17  Cloud9am             29614 non-null  float64
```

```

18 Cloud3pm      29176 non-null float64
19 Temp9am       51553 non-null float64
20 Temp3pm       50906 non-null float64
21 RainToday     50766 non-null object
22 RainTomorrow  52062 non-null object
dtypes: float64(16), object(7)
memory usage: 9.1+ MB

```

```
[ ]: test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29093 entries, 0 to 29092
Data columns (total 23 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date            29093 non-null  object
1   Location        29093 non-null  object
2   MinTemp         28922 non-null  float64
3   MaxTemp         28954 non-null  float64
4   Rainfall        28689 non-null  float64
5   Evaporation     16976 non-null  float64
6   Sunshine        16341 non-null  float64
7   WindGustDir     28580 non-null  object
8   WindGustSpeed   28617 non-null  float64
9   WindDir9am      27638 non-null  object
10  WindDir3pm      28783 non-null  object
11  WindSpeed9am    28938 non-null  float64
12  WindSpeed3pm    28932 non-null  float64
13  Humidity9am     28863 non-null  float64
14  Humidity3pm     28110 non-null  float64
15  Pressure9am     24890 non-null  float64
16  Pressure3pm     24889 non-null  float64
17  Cloud9am        17456 non-null  float64
18  Cloud3pm        16785 non-null  float64
19  Temp9am         28961 non-null  float64
20  Temp3pm         28278 non-null  float64
21  RainToday       28689 non-null  object
22  RainTomorrow    28685 non-null  object
dtypes: float64(16), object(7)
memory usage: 5.1+ MB

```

All the preprocessing steps which we are doing will also be done to the test dataset. This is why, we'll take both of the datasets together.

3.1.1 (a)

We can easily convert this to 0 or 1 using `pandas apply` method.

```
[ ]: train["RainTomorrow"] = train["RainTomorrow"].apply(lambda x: 0 if x=="No" else 1)
test["RainTomorrow"] = test["RainTomorrow"].apply(lambda x: 0 if x=="No" else 1)
```

```
[ ]: train.head()
```

```
[ ]:
      Date Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  \
0  01-12-2008  Albury    13.4    22.9      0.6          NaN        NaN
1  02-12-2008  Albury     7.4    25.1      0.0          NaN        NaN
2  03-12-2008  Albury    12.9    25.7      0.0          NaN        NaN
3  04-12-2008  Albury     9.2    28.0      0.0          NaN        NaN
4  05-12-2008  Albury    17.5    32.3      1.0          NaN        NaN

      WindGustDir  WindGustSpeed  WindDir9am  ...  Humidity9am  Humidity3pm  \
0              W           44.0           W  ...        71.0         22.0
1            WNW           44.0          NNW  ...        44.0         25.0
2            WSW           46.0           W  ...        38.0         30.0
3             NE           24.0           SE  ...        45.0         16.0
4             W           41.0          ENE  ...        82.0         33.0

      Pressure9am  Pressure3pm  Cloud9am  Cloud3pm  Temp9am  Temp3pm  RainToday  \
0         1007.7         1007.1        8.0      NaN     16.9     21.8         No
1         1010.6         1007.8        NaN      NaN     17.2     24.3         No
2         1007.6         1008.7        NaN      2.0     21.0     23.2         No
3         1017.6         1012.8        NaN      NaN     18.1     26.5         No
4         1010.8         1006.0        7.0      8.0     17.8     29.7         No

      RainTomorrow
0              0
1              0
2              0
3              0
4              0
```

[5 rows x 23 columns]

3.1.2 (b)

```
[ ]: object_cols_train = train.columns[train.dtypes == "object"]
object_cols_test = test.columns[test.dtypes == "object"]

object_cols_train, object_cols_test
```

```
[ ]: (Index(['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm',
            'RainToday'],
          dtype='object'),
      Index(['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm',
```

```
'RainToday'],
dtype='object'))
```

These are the object columns. We'll drop them:

```
[ ]: train.drop(object_cols_train, axis=1, inplace=True)
test.drop(object_cols_test, axis=1, inplace=True)
```

```
[ ]: train.head()
```

```
[ ]:
   MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  \
0    13.4    22.9      0.6          NaN        NaN          44.0
1     7.4    25.1      0.0          NaN        NaN          44.0
2    12.9    25.7      0.0          NaN        NaN          46.0
3     9.2    28.0      0.0          NaN        NaN          24.0
4    17.5    32.3      1.0          NaN        NaN          41.0

   WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm  Pressure9am  \
0          20.0          24.0          71.0          22.0        1007.7
1           4.0          22.0          44.0          25.0        1010.6
2          19.0          26.0          38.0          30.0        1007.6
3          11.0           9.0          45.0          16.0        1017.6
4           7.0          20.0          82.0          33.0        1010.8

   Pressure3pm  Cloud9am  Cloud3pm  Temp9am  Temp3pm  RainTomorrow
0        1007.1      8.0      NaN      16.9      21.8           0
1        1007.8      NaN      NaN      17.2      24.3           0
2        1008.7      NaN      2.0      21.0      23.2           0
3        1012.8      NaN      NaN      18.1      26.5           0
4        1006.0      7.0      8.0      17.8      29.7           0
```

3.1.3 (c)

Let's print some of the cells which are Null:

```
[ ]: train[train.isna()]
```

```
[ ]:
   MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  \
0        NaN      NaN      NaN          NaN        NaN          NaN
1        NaN      NaN      NaN          NaN        NaN          NaN
2        NaN      NaN      NaN          NaN        NaN          NaN
3        NaN      NaN      NaN          NaN        NaN          NaN
4        NaN      NaN      NaN          NaN        NaN          NaN
...      ...      ...      ...      ...      ...      ...
52057     NaN      NaN      NaN          NaN        NaN          NaN
52058     NaN      NaN      NaN          NaN        NaN          NaN
52059     NaN      NaN      NaN          NaN        NaN          NaN
52060     NaN      NaN      NaN          NaN        NaN          NaN
```

52061	NaN	NaN	NaN	NaN	NaN	NaN
	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	
...	
52057	NaN	NaN	NaN	NaN	NaN	
52058	NaN	NaN	NaN	NaN	NaN	
52059	NaN	NaN	NaN	NaN	NaN	
52060	NaN	NaN	NaN	NaN	NaN	
52061	NaN	NaN	NaN	NaN	NaN	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainTomorrow
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN
...	
52057	NaN	NaN	NaN	NaN	NaN	NaN
52058	NaN	NaN	NaN	NaN	NaN	NaN
52059	NaN	NaN	NaN	NaN	NaN	NaN
52060	NaN	NaN	NaN	NaN	NaN	NaN
52061	NaN	NaN	NaN	NaN	NaN	NaN

[52062 rows x 17 columns]

These are all the cells with null values. We'll fill them by means.

```
[ ]: cols = train.columns

for col in cols:
    train[col] = train[col].fillna(train[col].mean())
    test[col] = test[col].fillna(test[col].mean())
```

```
[ ]: train.isna().sum()
```

```
[ ]: MinTemp      0
      MaxTemp      0
      Rainfall     0
      Evaporation  0
      Sunshine     0
      WindGustSpeed 0
      WindSpeed9am 0
```



```

WindSpeed3pm      0
Humidity9am       0
Humidity3pm       0
Pressure9am       0
Pressure3pm       0
Cloud9am          0
Cloud3pm          0
Temp9am           0
Temp3pm           0
RainTomorrow      0
dtype: int64

```

```
[ ]: test.isna().sum()
```

```

[ ]: MinTemp      0
     MaxTemp      0
     Rainfall     0
     Evaporation  0
     Sunshine     0
     WindGustSpeed 0
     WindSpeed9am 0
     WindSpeed3pm 0
     Humidity9am   0
     Humidity3pm   0
     Pressure9am   0
     Pressure3pm   0
     Cloud9am      0
     Cloud3pm      0
     Temp9am       0
     Temp3pm       0
     RainTomorrow  0
dtype: int64

```

We see that there are no null values!

3.1.4 (d)

Let's make sure that all the columns are numerical:

```
[ ]: train.dtypes
```

```

[ ]: MinTemp      float64
     MaxTemp      float64
     Rainfall     float64
     Evaporation  float64
     Sunshine     float64
     WindGustSpeed float64
     WindSpeed9am float64

```

```

WindSpeed3pm      float64
Humidity9am       float64
Humidity3pm       float64
Pressure9am       float64
Pressure3pm       float64
Cloud9am          float64
Cloud3pm          float64
Temp9am           float64
Temp3pm           float64
RainTomorrow      int64
dtype: object

```

They are! Next, let's normalize the columns by:

$$X_{norm} = \frac{X - \min(X)}{\max(X) - \min(X)}$$

```

[ ]: def scale_one(col, df):
      minimum = df[col].min()
      maximum = df[col].max()
      return (df[col]-minimum)/(maximum-minimum)

```

```

[ ]: cols = train.columns

for col in cols:
    train[col] = scale_one(col, train)
    test[col] = scale_one(col, test)

```

```

[ ]: for col in cols:
      print(col, "Max", max(train[col]), max(test[col]))
      print(col, "Min", min(train[col]), min(test[col]))

```

```

MinTemp Max 1.0 1.0
MinTemp Min 0.0 0.0
MaxTemp Max 1.0 1.0
MaxTemp Min 0.0 0.0
Rainfall Max 1.0 1.0
Rainfall Min 0.0 0.0
Evaporation Max 1.0 1.0
Evaporation Min 0.0 0.0
Sunshine Max 1.0 1.0
Sunshine Min 0.0 0.0
WindGustSpeed Max 1.0 1.0
WindGustSpeed Min 0.0 0.0
WindSpeed9am Max 1.0 1.0
WindSpeed9am Min 0.0 0.0
WindSpeed3pm Max 1.0 1.0

```

```

WindSpeed3pm Min 0.0 0.0
Humidity9am Max 1.0 1.0
Humidity9am Min 0.0 0.0
Humidity3pm Max 1.0 1.0
Humidity3pm Min 0.0 0.0
Pressure9am Max 1.0 1.0
Pressure9am Min 0.0 0.0
Pressure3pm Max 1.0 1.0
Pressure3pm Min 0.0 0.0
Cloud9am Max 1.0 1.0
Cloud9am Min 0.0 0.0
Cloud3pm Max 1.0 1.0
Cloud3pm Min 0.0 0.0
Temp9am Max 1.0 1.0
Temp9am Min 0.0 0.0
Temp3pm Max 1.0 1.0
Temp3pm Min 0.0 0.0
RainTomorrow Max 1.0 1.0
RainTomorrow Min 0.0 0.0
Temp3pm Max 1.0 1.0
Temp3pm Min 0.0 0.0
RainTomorrow Max 1.0 1.0
RainTomorrow Min 0.0 0.0

```

Now, all values are between 0 and 1.

```
[ ]: train.head()
```

```

[ ]:      MinTemp    MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  \
0  0.569921  0.454139  0.001617    0.038129  0.536788    0.289062
1  0.411609  0.503356  0.000000    0.038129  0.536788    0.289062
2  0.556728  0.516779  0.000000    0.038129  0.536788    0.304688
3  0.459103  0.568233  0.000000    0.038129  0.536788    0.132812
4  0.678100  0.664430  0.002695    0.038129  0.536788    0.265625

      WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm  Pressure9am  \
0      0.153846      0.289157    0.701031    0.212121    0.452579
1      0.030769      0.265060    0.422680    0.242424    0.500832
2      0.146154      0.313253    0.360825    0.292929    0.450915
3      0.084615      0.108434    0.432990    0.151515    0.617304
4      0.053846      0.240964    0.814433    0.323232    0.504160

      Pressure3pm  Cloud9am  Cloud3pm  Temp9am  Temp3pm  RainTomorrow
0      0.477080  0.888889  0.560902  0.490196  0.439189          0.0
1      0.488964  0.492462  0.560902  0.497549  0.495495          0.0
2      0.504244  0.492462  0.250000  0.590686  0.470721          0.0
3      0.573854  0.492462  0.560902  0.519608  0.545045          0.0
4      0.458404  0.777778  1.000000  0.512255  0.617117          0.0

```

```
[ ]: test.head()
```

```
[ ]:      MinTemp    MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  \
0  0.664835  0.954198      0.0      0.121439  0.922535      0.387097
1  1.000000  0.949109      0.0      0.121439  0.859155      0.387097
2  0.673077  0.592875      0.0      0.121439  0.661972      0.322581
3  0.527473  0.577608      0.0      0.121439  0.922535      0.306452
4  0.491758  0.531807      0.0      0.121439  0.873239      0.354839

      WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm  Pressure9am  \
0      0.412698      0.238095      0.262626      0.090909      0.567753
1      0.492063      0.476190      0.090909      0.212121      0.437393
2      0.349206      0.444444      0.626263      0.424242      0.492281
3      0.269841      0.476190      0.424242      0.282828      0.514580
4      0.174603      0.492063      0.474747      0.333333      0.449400

      Pressure3pm  Cloud9am  Cloud3pm  Temp9am  Temp3pm  RainTomorrow
0      0.5216  0.472588  0.491987  0.788235  0.955882      0.0
1      0.4912  0.472588  0.491987  1.000000  0.781863      0.0
2      0.5184  0.875000  0.250000  0.600000  0.580882      0.0
3      0.5184  0.000000  0.000000  0.595294  0.615196      0.0
4      0.5008  0.125000  0.375000  0.604706  0.534314      0.0
```

3.2 Problem 3.2

We have given the the formal definition of the optimization problem in the section 2.1. Here, we'll just optimize the cost function using `scipy.optimize.minimize` function.

3.2.1 Defining Some Functions

First, let's define some functions

```
[ ]: def sigmoid(x):
      """The sigmoid function"""
      return 1 / (1 + np.exp(-x))

def _get_loss(y_hat, y_true):
      """The loss function for logistic regression"""
      m = len(y_hat)
      return np.sum(-y_true * np.log(y_hat) - (1 - y_true) * (np.log(1 - y_hat)))
      ↪ / m

def _get_yhat(X, w):
      """To get the predicted values for a given weight vector w"""
      return sigmoid(np.dot(X, w.T))

def _get_weights(n):
      """Gets a random weight vector of size n"""
```

```
w = np.random.random((n))
return w
```

```
[ ]: def loss(w, X, y_true):
    """
    The objective function. This is the function which will be used to optimize
    """
    y_hat = sigmoid(np.dot(X, w.T))
    m = len(y_hat)
    loss = np.sum(-y_true * np.log(y_hat) - (1 - y_true) * (np.log(1 - y_hat)))
    ↪ / m
    return loss
```

3.2.2 Creating X and y

Here, we'll separate X and y . Also, as our hypothesis function contains a bias term, we'll add a vector of 1 to the X so that b is contained in w .

```
[ ]: X = train.drop("RainTomorrow", axis=1).values
y = train["RainTomorrow"]
ones = np.ones((X.shape[0], 1))
X = np.append(ones, X, axis=1)

test_X = test.drop("RainTomorrow", axis=1).values
test_y = test["RainTomorrow"]
ones = np.ones((test_X.shape[0], 1))
test_X = np.append(ones, test_X, axis=1)
```

We won't be creating a train-test split as we already have a test set!

3.2.3 Initializing w

Next step is to initialize the weights. Instead of initializing them by zeros or ones, I'm doing a random initialization.

```
[ ]: np.random.seed(42)
m, n = X.shape
init_weights = _get_weights(n)
init_weights
```

```
[ ]: array([0.37454012, 0.95071431, 0.73199394, 0.59865848, 0.15601864,
          0.15599452, 0.05808361, 0.86617615, 0.60111501, 0.70807258,
          0.02058449, 0.96990985, 0.83244264, 0.21233911, 0.18182497,
          0.18340451, 0.30424224])
```

3.2.4 Optimization

Before optimization, we'll create a callback function which will be used for plotting purposes.

```
[ ]: train_losses = []
      test_losses = []
      epoch = 1
      def callback(w):
          global epoch
          print(f"On Iteration {epoch:3.0f}", end="\r")
          train_loss = loss(w, X, y)
          test_loss = loss(w, test_X, test_y)
          train_losses.append(train_loss)
          test_losses.append(test_loss)
          epoch+=1
```

The problem specifies that we need to use Truncated Newton's method. We'll use `scipy.optimize.minimize` function for this. We'll use TNC method for this.

```
[ ]: train_losses = []
      test_losses = []

      optimizer = minimize(
          loss,
          init_weights,
          args=(X, y),
          method="TNC",
          options={"maxiter": 1000},
          callback=callback,
          tol=1e-8
      )
```

On Iteration 142

```
[ ]: w_final = optimizer["x"]
      w_final
```

```
[ ]: array([-5.57841148,  0.3942323 ,  1.43262747,  2.35296417, -0.07447357,
          -1.57413976,  5.29430146,  0.65212976, -1.73175333,  1.87786876,
           4.16884424,  0.41801538, -1.55171722,  0.11354733,  0.90514991,
           0.5589355 , -1.36817804])
```

So, these are the weights. Let's see the accuracy of the model on the train dataset to see whether the model is working or not.

```
[ ]: def accuracy(y_true, y_pred):
      return np.mean(y_true==y_pred)
```

```
[ ]: y_pred = _get_yhat(X, w_final)
      y_pred = np.round(y_pred)
```

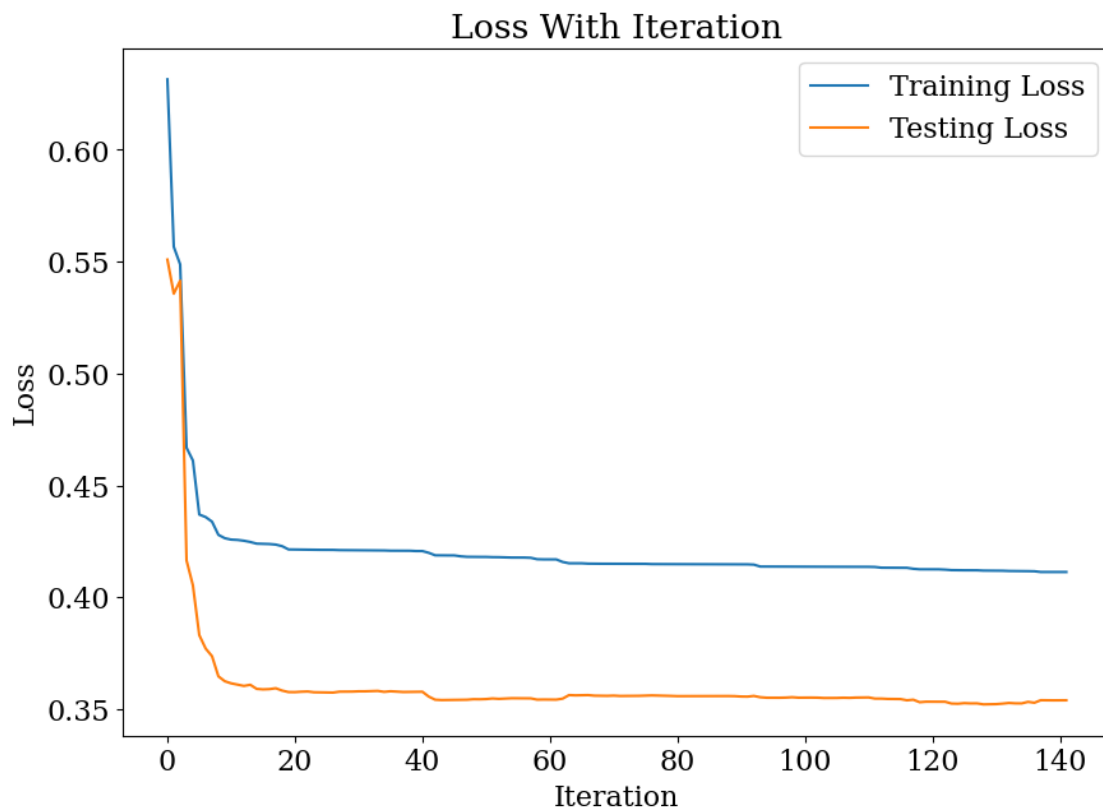
```
[ ]: train_acc = accuracy(y, y_pred)
      print(f"Training accuracy is: {train_acc*100:.2f}%")
```

Training accuracy is: 82.03%

3.3 Problem 3.3

I've already used callback function to store the loss function at each iteration. Let's plot them!

```
[ ]: fig, ax = plt.subplots()
      ax.plot(train_losses, label = "Training Loss")
      ax.plot(test_losses, label = "Testing Loss")
      ax.set_xlabel("Iteration")
      ax.set_ylabel("Loss")
      ax.set_title("Loss With Iteration")
      ax.legend()
      fig.savefig(os.path.join(DATA_DIR, "0301.png"))
```



Okay, we see that both the training and testing accuracy is decreasing with iteration. No significant change occurs after about 20 iterations.

3.4 Problem 3.4

We have the model as `optimizer`. We also have the weights of the model which is all we need. Let's determine the testing accuracy.

```
[ ]: y_pred = _get_yhat(test_X, w_final)
      y_pred = np.round(y_pred)

      test_acc = accuracy(test_y, y_pred)
      print(f"Testing accuracy is: {test_acc*100:.2f}%")
```

Testing accuracy is: 84.78%

So, testing accuracy is greater than the train accuracy. Means model is not overfitting.

4 Extra

I implemented logistic regression using gradient descent along with linear regression. The class `LogisticGradientDescent` is a subclass of `BatchGradientDescent`. Let's see how this model performs.

```
[ ]: from GD import LogisticGradientDescent

      lgd = LogisticGradientDescent(fit_intercept=True, tol = 1e-6)
```

```
[ ]: lgd.fit(X, y, epochs=10000, learning_rate=0.1, verbose=0)
```

10000/10000 [=====] 100.0%

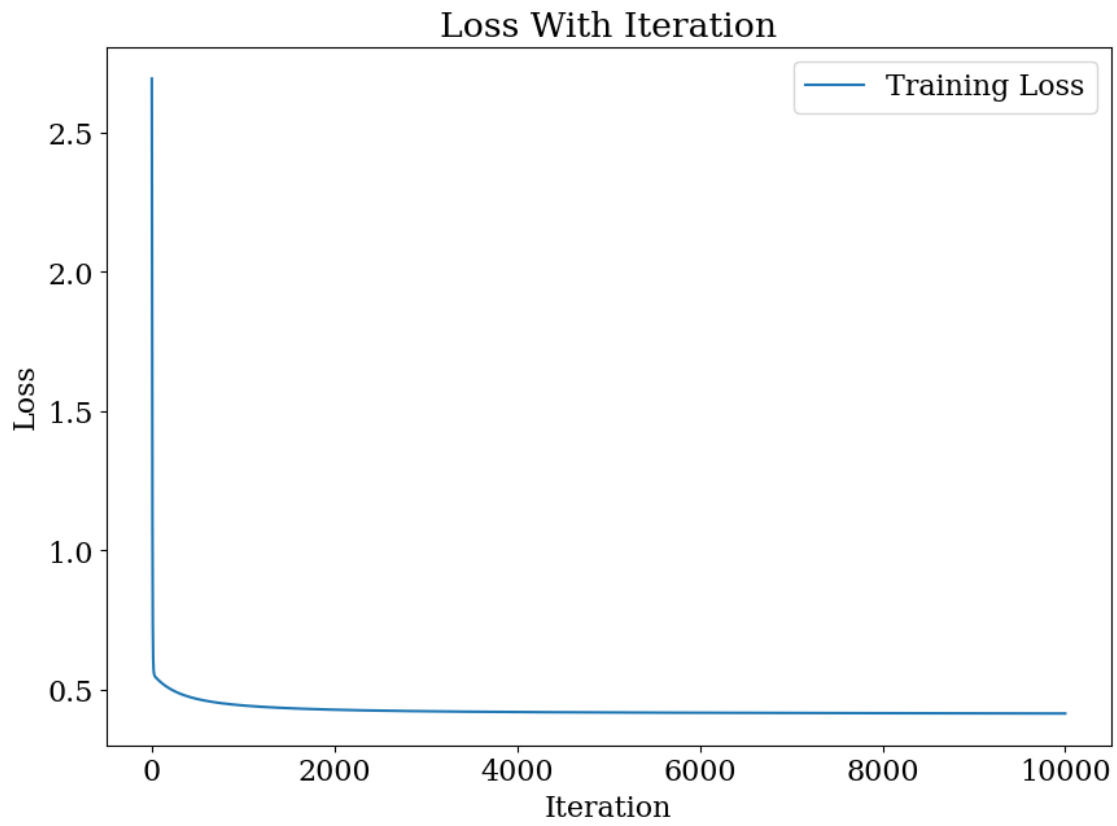
```
[ ]: train_acc = lgd.score(X, y)
      test_acc = lgd.score(test_X, test_y)
      print(f"Training accuracy is: {train_acc*100:.2f}%")
      print(f"Testing accuracy is: {test_acc*100:.2f}%")
```

Training accuracy is: 81.88%

Testing accuracy is: 84.70%

The accuracy is almost same as that by optimizing. Let's plot the cost function.

```
[ ]: fig, ax = plt.subplots()
      ax.plot(lgd._losses, label = "Training Loss")
      ax.set_xlabel("Iteration")
      ax.set_ylabel("Loss")
      ax.set_title("Loss With Iteration")
      ax.legend()
      fig.savefig(os.path.join(DATA_DIR, "0302.png"))
```

The curve is much smoother.