

Assignment_5

March 12, 2023

1 Imports

```
[ ]: from nnet.model import Sequential
from nnet.layers import (
    Input,
    Dense,
    Conv2D,
    MaxPool2D,
    Flatten,
)
from nnet.utils import one_hot
import numpy as np
import os
from scipy.io import loadmat
import matplotlib.pyplot as plt
```

2 Loading Data

```
[ ]: DATA_DIR = "data"
X_train = loadmat(os.path.join(DATA_DIR, "train_images.mat"))["train_images"]
X_test = loadmat(os.path.join(DATA_DIR, "test_images.mat"))["test_images"]
y_train = loadmat(os.path.join(DATA_DIR, "train_labels.mat"))["train_labels"]
y_test = loadmat(os.path.join(DATA_DIR, "test_labels.mat"))["test_labels"]
```

2.1 Preprocessing

```
[ ]: # Making the data compatible with the model
X_train = X_train.T
X_test = X_test.T
X_train = X_train[:, :, np.newaxis, :]
X_test = X_test[:, :, np.newaxis, :]
X_train.shape, X_test.shape
```

```
[ ]: ((28, 28, 1, 1000), (28, 28, 1, 1000))
```

```
[ ]: # Normalizing the data
X_train = X_train / X_train.max() - 0.5
X_test = X_test / X_test.max() - 0.5
```

```
[ ]: # One hot encoding the labels
y_train = np.squeeze(y_train.T)
y_test = np.squeeze(y_test.T)
y_train = one_hot(y_train, 10)
y_test = one_hot(y_test, 10)

y_train.shape, y_test.shape
```

```
[ ]: ((10, 1000), (10, 1000))
```

3 Problems

3.1 Problem 1

Note: I'm using the updated version of the module which I created for the previous assignment. To see how to work with the module, please refer to the previous assignment.

Here is the model:

```
[ ]: model = Sequential("CNN Model")
inp = Input((28, 28, 1), name="Input")
conv = Conv2D(9, (3, 3), padding="valid", activation="relu", name="Conv")
pool = MaxPool2D((2, 2), 2, name="MaxPool")
flat = Flatten(name="Flatten")
dense = Dense(10, name="Output", activation="softmax")

model.add(inp)
model.add(conv)
model.add(pool)
model.add(flat)
model.add(dense)

model.summary()
```

Model: CNN Model

Name	Input Shapes	Output Shapes	Weight Shapes	Bias Shapes
# Parameters				

Input 0	(28, 28, 1)	(28, 28, 1)	None	None
Conv 90	(28, 28, 1)	(13, 13, 9)	(3, 3, 1, 9)	(9, 1)
MaxPool 0	(13, 13, 9)	(6, 6, 9)	None	None
Flatten 0	(6, 6, 9)	(324,)	None	None
Output 3250	(324,)	(10,)	(10, 324)	(10, 1)

```
=====
=====
Total Parameters: 3340
-----
-----
```

So, the model has just 3340 parameters. Would this be enough to learn the data? Let's find out. For this, we compile the model:

```
[ ]: model.compile(
    loss="categorical_crossentropy",
    metrics=["accuracy"],
    initializer="glorot",
)
```

Finally, fit the model:

```
[ ]: history = model.fit(X_train, y_train, epochs=50, batch_size=32, lr=0.05,
    ↪ verbose=10)
```

```
Epoch 0001/0050 | Loss: 2.56050 | Accuracy: 0.18200 |
Epoch 0002/0050 | Loss: 2.44336 | Accuracy: 0.20800 |
Epoch 0003/0050 | Loss: 2.35812 | Accuracy: 0.21400 |
Epoch 0004/0050 | Loss: 2.29177 | Accuracy: 0.21200 |
Epoch 0005/0050 | Loss: 2.23784 | Accuracy: 0.21100 |
Epoch 0006/0050 | Loss: 2.19254 | Accuracy: 0.21200 |
Epoch 0007/0050 | Loss: 2.15412 | Accuracy: 0.30300 |
Epoch 0008/0050 | Loss: 2.12118 | Accuracy: 0.30900 |
Epoch 0009/0050 | Loss: 2.09269 | Accuracy: 0.31700 |
```

Epoch 0010/0050	Loss: 2.06748	Accuracy: 0.32400	
Epoch 0011/0050	Loss: 2.04555	Accuracy: 0.32700	
Epoch 0012/0050	Loss: 2.02632	Accuracy: 0.32900	
Epoch 0013/0050	Loss: 2.00957	Accuracy: 0.33200	
Epoch 0014/0050	Loss: 1.99492	Accuracy: 0.33000	
Epoch 0015/0050	Loss: 1.98208	Accuracy: 0.32900	
Epoch 0016/0050	Loss: 1.97072	Accuracy: 0.33200	
Epoch 0017/0050	Loss: 1.96073	Accuracy: 0.33300	
Epoch 0018/0050	Loss: 1.95208	Accuracy: 0.33000	
Epoch 0019/0050	Loss: 1.94455	Accuracy: 0.33000	
Epoch 0020/0050	Loss: 1.93779	Accuracy: 0.32900	
Epoch 0021/0050	Loss: 1.93199	Accuracy: 0.33200	
Epoch 0022/0050	Loss: 1.92699	Accuracy: 0.33500	
Epoch 0023/0050	Loss: 1.92250	Accuracy: 0.33300	
Epoch 0024/0050	Loss: 1.91830	Accuracy: 0.33600	
Epoch 0025/0050	Loss: 1.91448	Accuracy: 0.33700	
Epoch 0026/0050	Loss: 1.91109	Accuracy: 0.33600	
Epoch 0027/0050	Loss: 1.90838	Accuracy: 0.33900	
Epoch 0028/0050	Loss: 1.90624	Accuracy: 0.34300	
Epoch 0029/0050	Loss: 1.90439	Accuracy: 0.34700	
Epoch 0030/0050	Loss: 1.90300	Accuracy: 0.34900	
Epoch 0031/0050	Loss: 1.90178	Accuracy: 0.35000	
Epoch 0032/0050	Loss: 1.90051	Accuracy: 0.34900	
Epoch 0033/0050	Loss: 1.89907	Accuracy: 0.35000	
Epoch 0034/0050	Loss: 1.89647	Accuracy: 0.34700	
Epoch 0035/0050	Loss: 1.89163	Accuracy: 0.34800	
Epoch 0036/0050	Loss: 1.88653	Accuracy: 0.35200	
Epoch 0037/0050	Loss: 1.88185	Accuracy: 0.35200	
Epoch 0038/0050	Loss: 1.87717	Accuracy: 0.35400	
Epoch 0039/0050	Loss: 1.87196	Accuracy: 0.35500	
Epoch 0040/0050	Loss: 1.86731	Accuracy: 0.35600	
Epoch 0041/0050	Loss: 1.86298	Accuracy: 0.35800	
Epoch 0042/0050	Loss: 1.85855	Accuracy: 0.35800	
Epoch 0043/0050	Loss: 1.85395	Accuracy: 0.35700	
Epoch 0044/0050	Loss: 1.84980	Accuracy: 0.36000	
Epoch 0045/0050	Loss: 1.84476	Accuracy: 0.36000	
Epoch 0046/0050	Loss: 1.83999	Accuracy: 0.36200	
Epoch 0047/0050	Loss: 1.83557	Accuracy: 0.36200	
Epoch 0048/0050	Loss: 1.83127	Accuracy: 0.36300	
Epoch 0049/0050	Loss: 1.82735	Accuracy: 0.36700	
Epoch 0050/0050	Loss: 1.82316	Accuracy: 0.36800	

The accuracy is very low. This is because the model implemented above is very simple. It has just 3340 parameters. This is not enough to learn the data.

```
[ ]: train_acc = model.evaluate(X_train, y_train, metric="accuracy")
      print(f"Train Accuracy: {train_acc*100:.2f}%")
```

Train Accuracy: 36.80%

So, the training accuracy is about 37%. Let's see how the model performs on the test data.

3.2 Problem 2

```
[ ]: test_acc = model.evaluate(X_test, y_test, metric="accuracy")
      print(f"Test Accuracy: {test_acc*100:.2f}%")
```

Test Accuracy: 34.30%

Okay, the test accuracy is very close to the training accuracy. This means that the model is not overfitting.